

Assignment 2 - Neural Networks

This notebook explores different configurations of a neural network for classifying IMDB movie reviews.

Step 1: Load and Preprocess Data

We will load the IMDB dataset and preprocess it for training.

```
In [18]: import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

# Load the IMDB dataset
from tensorflow.keras.datasets import imdb

num_words = 10000 # Restrict to top 10,000 words
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=num_words)

# Function to vectorize data (one-hot encoding)
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        results[i, sequence] = 1.
    return results

# Preprocess data
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
y_train = np.asarray(train_labels).astype('float32')
y_test = np.asarray(test_labels).astype('float32')

# Split training data into training and validation sets
```

```

x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

print(f"Training data shape: {x_train.shape}")
print(f"Validation data shape: {x_val.shape}")
print(f"Test data shape: {x_test.shape}")

```

Training data shape: (25000, 10000)
 Validation data shape: (10000, 10000)
 Test data shape: (25000, 10000)

Step 2: Define Function to Build Models

```

In [19]: # Define the model function
def build_model(hidden_layers=2, hidden_units=16, activation="relu", loss_fn="binary_crossentropy"):
    model = keras.Sequential()
    model.add(layers.Dense(hidden_units, activation=activation, input_shape=(10000,)))
    for _ in range(hidden_layers - 1):
        model.add(layers.Dense(hidden_units, activation=activation))
    model.add(layers.Dense(1, activation="sigmoid")) # Output Layer

    model.compile(optimizer="rmsprop", loss=loss_fn, metrics=["accuracy"])
    return model

```

Step 3: Experimenting with Different Architectures

We will test models with different numbers of hidden layers and units.

```

In [20]: # Define all combinations of hidden layers and units per layer
configurations = [(layers, units) for layers in range(1, 4) for units in [8, 16, 32, 64]]

# Dictionary to store results
config_results = {}

# Train models and evaluate
for layers_count, units in configurations:
    print(f"Training model with {layers_count} hidden layers and {units} units per layer...")

```

```

model = build_model(hidden_layers=layers_count, hidden_units=units)
history = model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                    validation_data=(x_val, y_val), verbose=0)
eval_results = model.evaluate(x_test, y_test, verbose=0) # [Loss, Accuracy]

config_results[(layers_count, units)] = eval_results

# Convert results to DataFrame
config_results_df = pd.DataFrame.from_dict(config_results, orient="index", columns=["Loss", "Accuracy"])
config_results_df.index.names = ["Hidden Layers", "Units"]
print(config_results_df)

```

Training model with 1 hidden layers and 8 units per layer...

C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

Training model with 1 hidden layers and 16 units per layer...
Training model with 1 hidden layers and 32 units per layer...
Training model with 1 hidden layers and 64 units per layer...
Training model with 2 hidden layers and 8 units per layer...
Training model with 2 hidden layers and 16 units per layer...
Training model with 2 hidden layers and 32 units per layer...
Training model with 2 hidden layers and 64 units per layer...
Training model with 3 hidden layers and 8 units per layer...
Training model with 3 hidden layers and 16 units per layer...
Training model with 3 hidden layers and 32 units per layer...
Training model with 3 hidden layers and 64 units per layer...

```

	Loss	Accuracy
Hidden Layers, Units		
(1, 8)	0.361747	0.87064
(1, 16)	0.404524	0.86548
(1, 32)	0.429601	0.86428
(1, 64)	0.490549	0.85756
(2, 8)	0.461372	0.86404
(2, 16)	0.627186	0.85000
(2, 32)	0.672112	0.86132
(2, 64)	0.634104	0.86564
(3, 8)	0.529646	0.85684
(3, 16)	0.705806	0.85688
(3, 32)	0.780766	0.85996
(3, 64)	0.763498	0.86708

Step 4: Visualizing Performance of Different Architectures

```

In [21]: # Plot Accuracy Comparison
plt.figure(figsize=(10, 6))
sns.barplot(x=[f"{layers}L-{units}U" for layers, units in config_results.keys()],
            y=[results[1] for results in config_results.values()], palette="viridis")

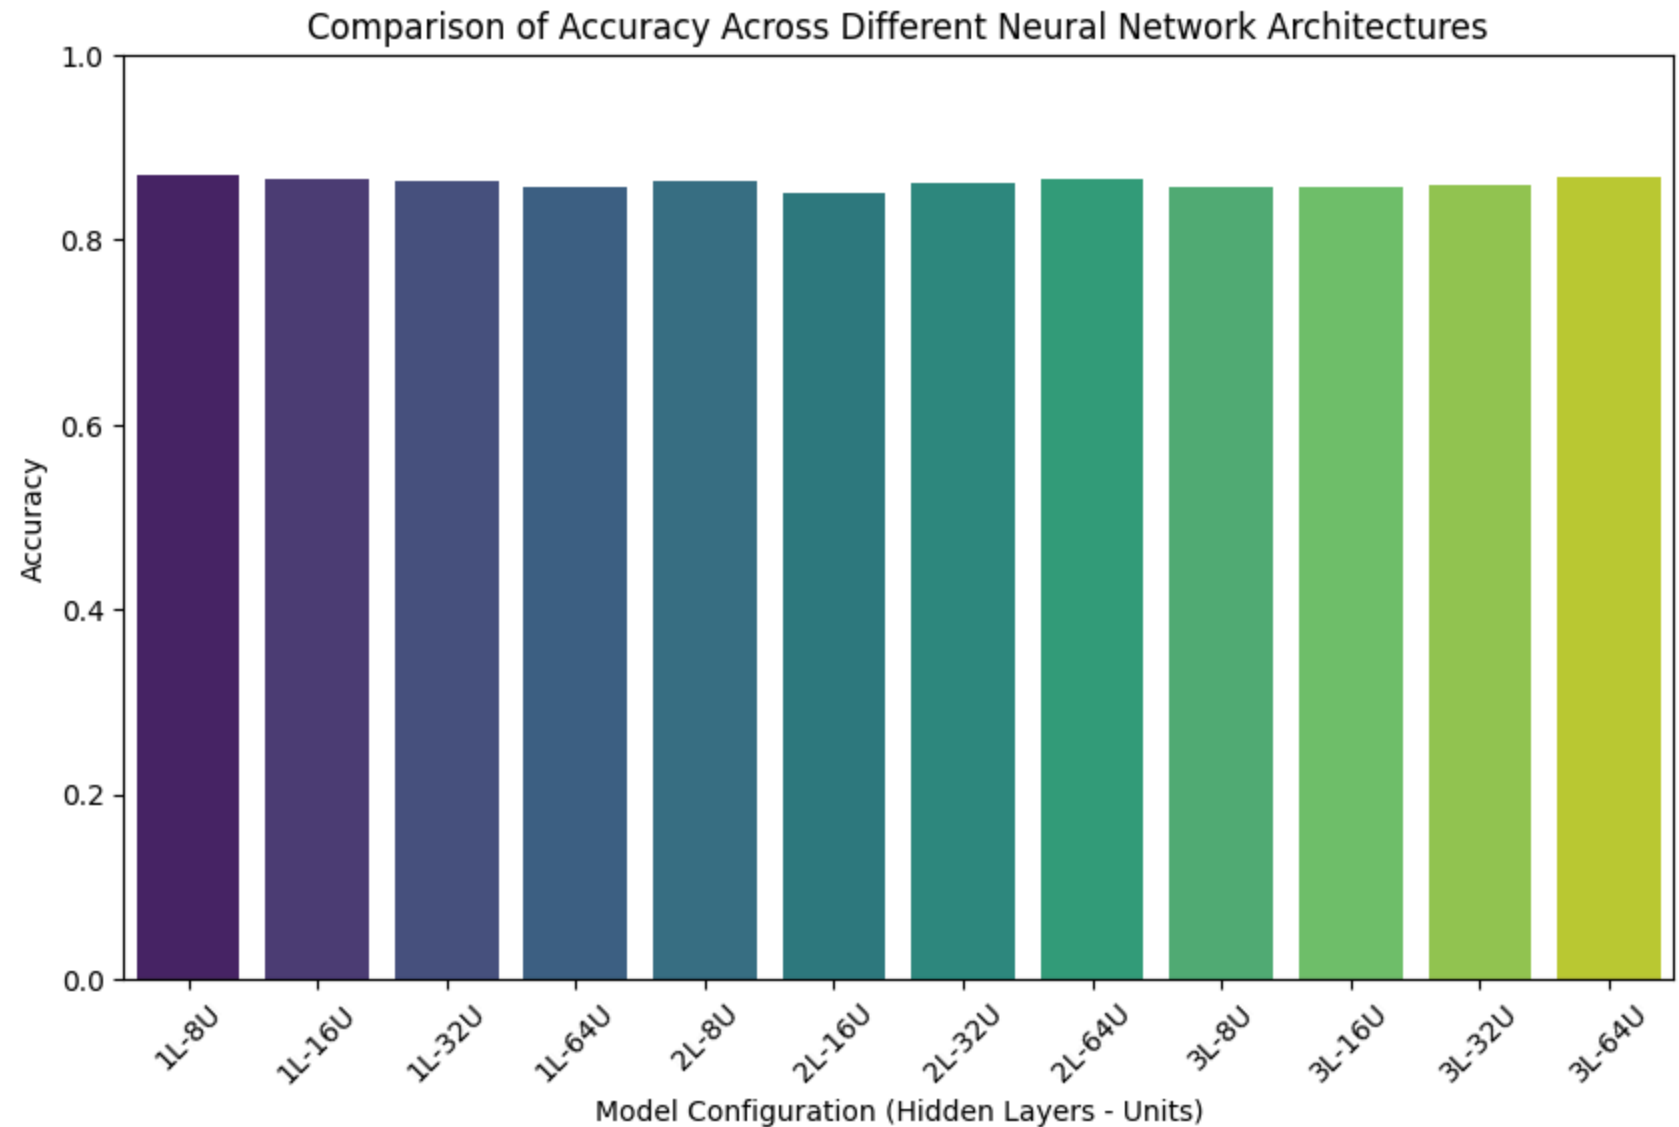
plt.xlabel("Model Configuration (Hidden Layers - Units)")
plt.ylabel("Accuracy")
plt.title("Comparison of Accuracy Across Different Neural Network Architectures")
plt.xticks(rotation=45)
plt.ylim(0, 1)
plt.show()

```

C:\Users\zhanguu\AppData\Local\Temp\ipykernel_26204\2984978426.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=[f"{layers}L-{units}U" for layers, units in config_results.keys()],
```



Step 5: Changing Loss Function to MSE

```
In [22]: # Dictionary to store MSE results
mse_results = {}

# Train models using MSE Loss function
for layers_count, units in configurations:
    print(f"Training model with {layers_count} hidden layers and {units} units per layer using MSE loss...")

    mse_model = build_model(hidden_layers=layers_count, hidden_units=units, loss_fn="mse")
    mse_history = mse_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                               validation_data=(x_val, y_val), verbose=0)
    eval_results = mse_model.evaluate(x_test, y_test, verbose=0) # [Loss, Accuracy]

    mse_results[(layers_count, units)] = eval_results

# Convert MSE results to DataFrame
mse_results_df = pd.DataFrame.from_dict(mse_results, orient="index", columns=["Loss", "Accuracy"])
mse_results_df.index.names = ["Hidden Layers", "Units"]

# Combine BCE and MSE results into one DataFrame
comparison_df = config_results_df.join(mse_results_df, lsuffix="_BCE", rsuffix="_MSE")
print(comparison_df)

# --- Plot Accuracy Comparison ---
plt.figure(figsize=(12, 6))
x_labels = [f"{layers}L-{units}U" for layers, units in configurations]

plt.plot(x_labels, comparison_df["Accuracy_BCE"], marker='o', label="BCE Accuracy", linestyle='--', color='blue')
plt.plot(x_labels, comparison_df["Accuracy_MSE"], marker='s', label="MSE Accuracy", linestyle='-', color='red')

plt.xlabel("Model Configuration (Hidden Layers - Units)")
plt.ylabel("Accuracy")
plt.title("Accuracy Comparison: Binary Crossentropy vs. MSE")
plt.xticks(rotation=45)
plt.legend()
plt.ylim(0, 1)
plt.show()

# --- Plot Loss Comparison ---
```

```

plt.figure(figsize=(12, 6))
plt.plot(x_labels, comparison_df["Loss_BCE"], marker='o', label="BCE Loss", linestyle='--', color='blue')
plt.plot(x_labels, comparison_df["Loss_MSE"], marker='s', label="MSE Loss", linestyle='--', color='red')

plt.xlabel("Model Configuration (Hidden Layers - Units)")
plt.ylabel("Loss")
plt.title("Loss Comparison: Binary Crossentropy vs. MSE")
plt.xticks(rotation=45)
plt.legend()
plt.show()

```

Training model with 1 hidden layers and 8 units per layer using MSE loss...

C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Training model with 1 hidden layers and 16 units per layer using MSE loss...

Training model with 1 hidden layers and 32 units per layer using MSE loss...

Training model with 1 hidden layers and 64 units per layer using MSE loss...

Training model with 2 hidden layers and 8 units per layer using MSE loss...

Training model with 2 hidden layers and 16 units per layer using MSE loss...

Training model with 2 hidden layers and 32 units per layer using MSE loss...

Training model with 2 hidden layers and 64 units per layer using MSE loss...

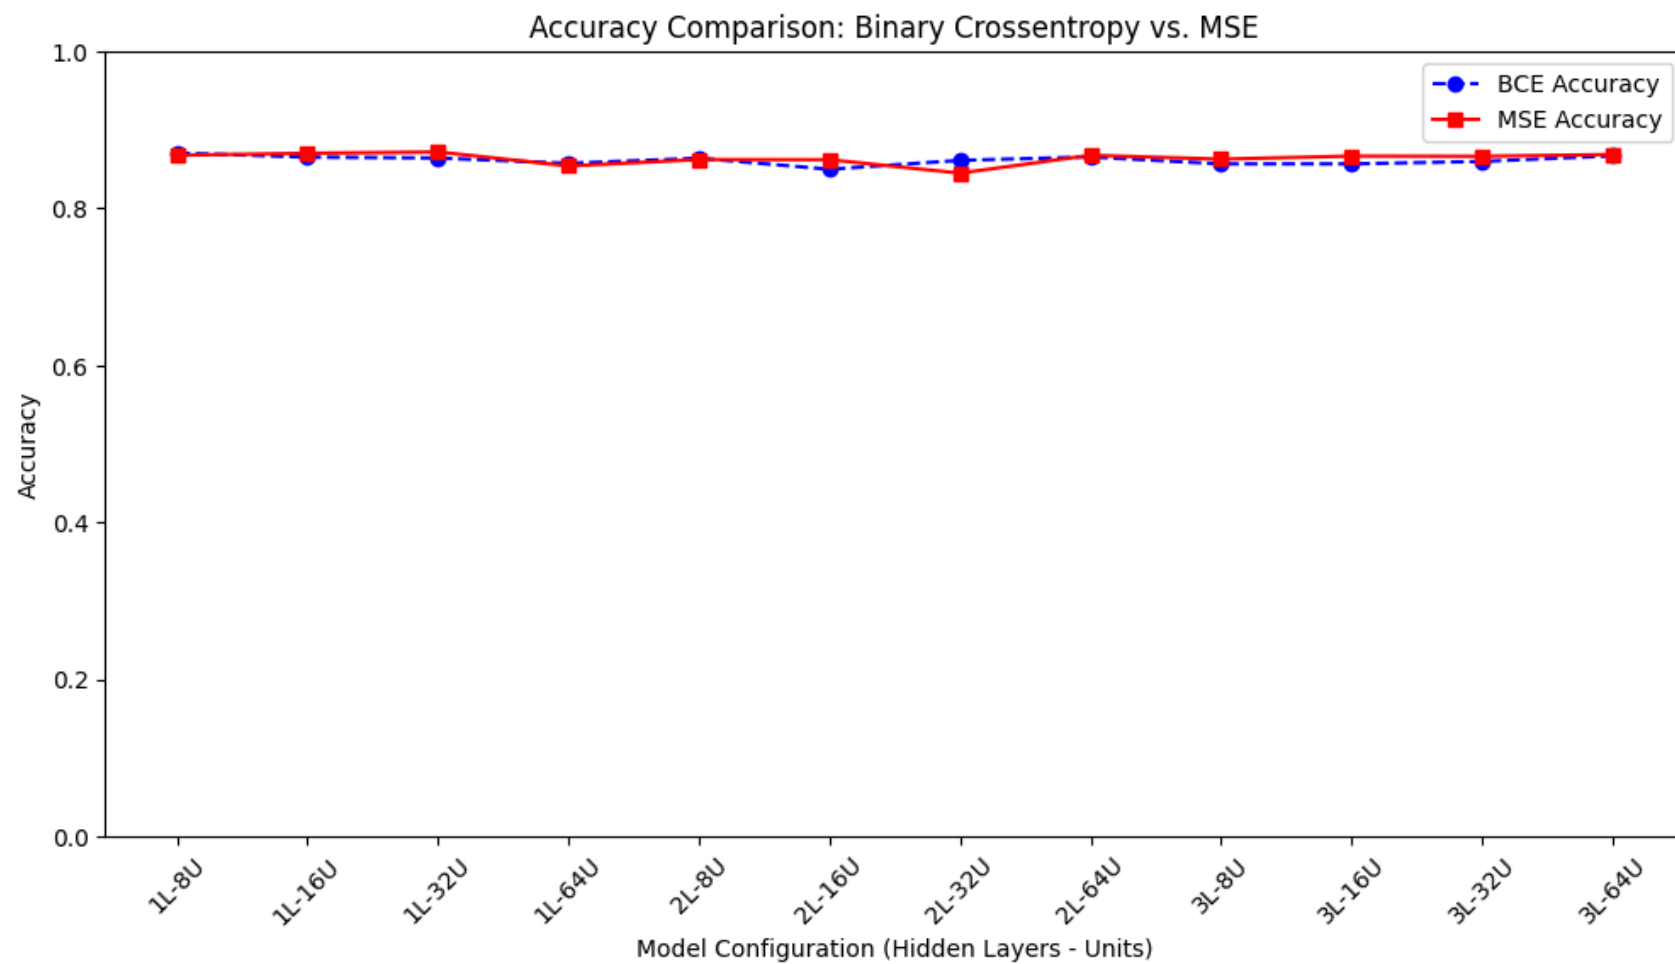
Training model with 3 hidden layers and 8 units per layer using MSE loss...

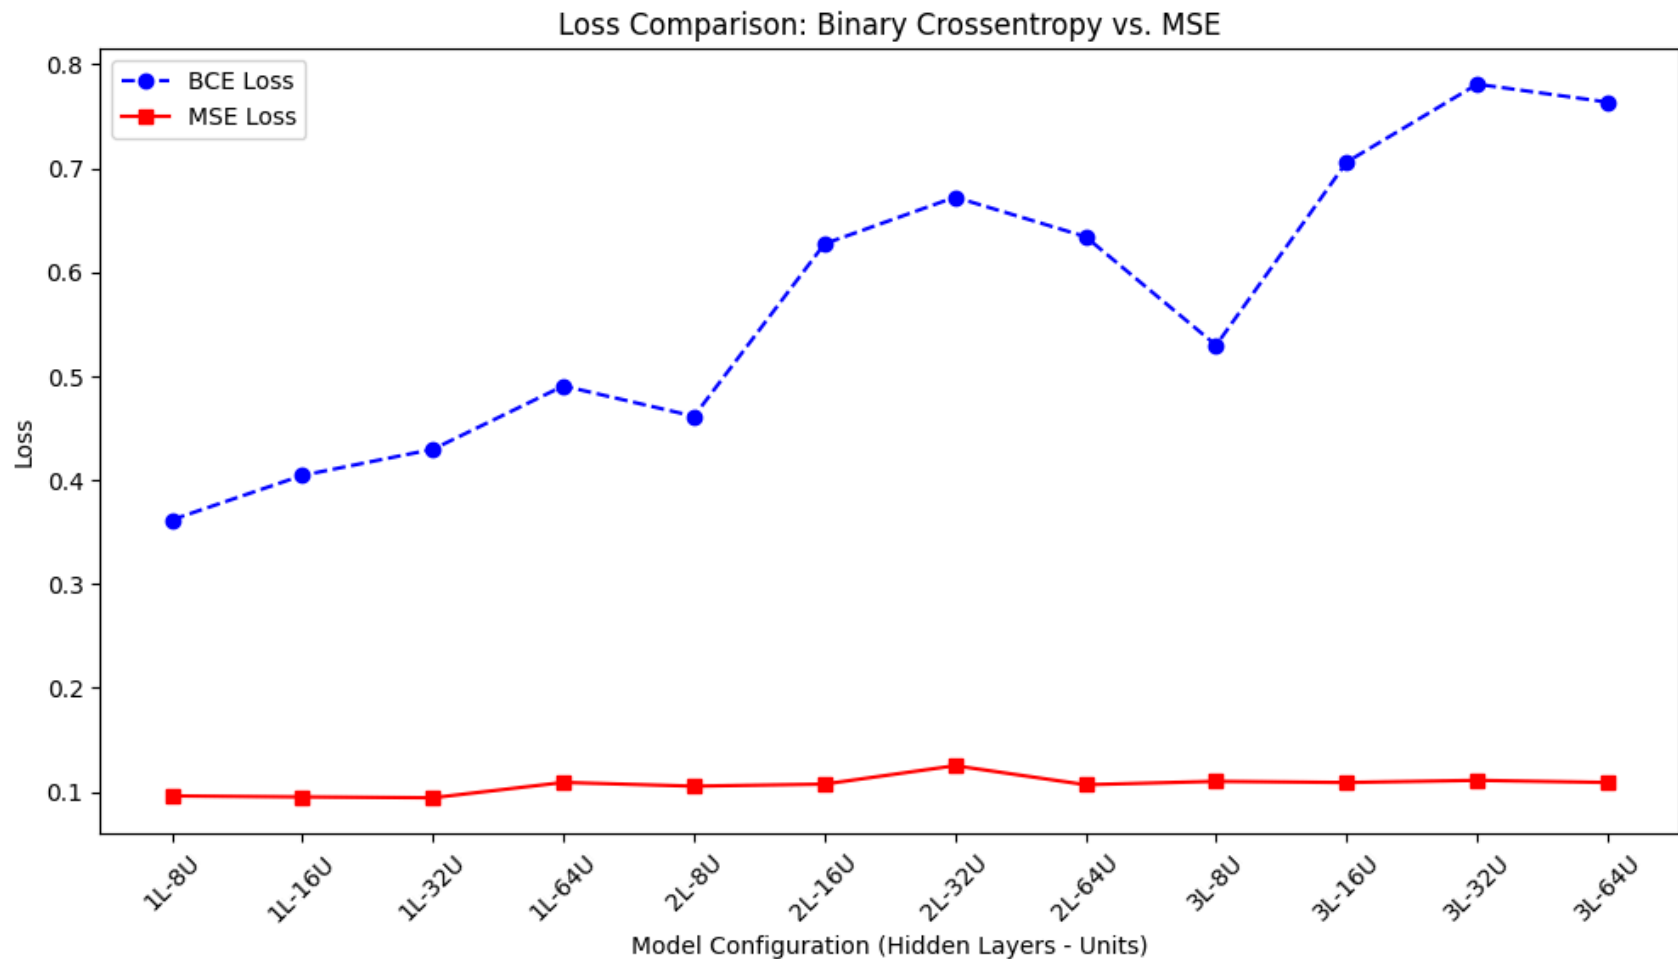
Training model with 3 hidden layers and 16 units per layer using MSE loss...

Training model with 3 hidden layers and 32 units per layer using MSE loss...

Training model with 3 hidden layers and 64 units per layer using MSE loss...

Hidden Layers, Units	Loss_BCE	Accuracy_BCE	Loss_MSE	Accuracy_MSE
(1, 8)	0.361747	0.87064	0.096216	0.86788
(1, 16)	0.404524	0.86548	0.095235	0.87036
(1, 32)	0.429601	0.86428	0.094448	0.87208
(1, 64)	0.490549	0.85756	0.109183	0.85416
(2, 8)	0.461372	0.86404	0.105588	0.86236
(2, 16)	0.627186	0.85000	0.107584	0.86216
(2, 32)	0.672112	0.86132	0.125297	0.84504
(2, 64)	0.634104	0.86564	0.106975	0.86784
(3, 8)	0.529646	0.85684	0.110227	0.86284
(3, 16)	0.705806	0.85688	0.109105	0.86684
(3, 32)	0.780766	0.85996	0.111202	0.86656
(3, 64)	0.763498	0.86708	0.109137	0.86880





```
In [24]: # Train a model with (2 hidden layers, 16 units) using BCE for 20 epochs
bce_model = build_model(hidden_layers=2, hidden_units=16, loss_fn="binary_crossentropy")
bce_history = bce_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                           validation_data=(x_val, y_val), verbose=1)

# Train a model with (2 hidden layers, 16 units) using MSE for 20 epochs
mse_model = build_model(hidden_layers=2, hidden_units=16, loss_fn="mse")
mse_history = mse_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                           validation_data=(x_val, y_val), verbose=1)

# Plot Training Evolution for Loss and Accuracy over all 20 Epochs
plt.figure(figsize=(12, 5))
```

```

# Plot Loss Evolution
plt.subplot(1, 2, 1)
plt.plot(range(1, 21), bce_history.history["loss"], label="BCE Loss", color="blue")
plt.plot(range(1, 21), mse_history.history["loss"], label="MSE Loss", color="red", linestyle="--")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss Evolution: BCE vs. MSE")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


# Plot Accuracy Evolution
plt.subplot(1, 2, 2)
plt.plot(range(1, 21), bce_history.history["accuracy"], label="BCE Accuracy", color="blue")
plt.plot(range(1, 21), mse_history.history["accuracy"], label="MSE Accuracy", color="red", linestyle="--")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Training Accuracy Evolution: BCE vs. MSE")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


plt.tight_layout()
plt.show()


```


C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/20
30/30  1s 22ms/step - accuracy: 0.6789 - loss: 0.6170 - val_accuracy: 0.8665 - val_loss: 0.4149


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8894 - loss: 0.3624 - val_accuracy: 0.8811 - val_loss: 0.3260


Epoch 3/20
30/30  0s 11ms/step - accuracy: 0.9173 - loss: 0.2575 - val_accuracy: 0.8901 - val_loss: 0.2857


Epoch 4/20
30/30  0s 11ms/step - accuracy: 0.9347 - loss: 0.2017 - val_accuracy: 0.8743 - val_loss: 0.3120


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.9441 - loss: 0.1706 - val_accuracy: 0.8806 - val_loss: 0.2992


Epoch 6/20
30/30  0s 12ms/step - accuracy: 0.9528 - loss: 0.1432 - val_accuracy: 0.8855 - val_loss: 0.2869


Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9651 - loss: 0.1181 - val_accuracy: 0.8849 - val_loss: 0.2950


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9719 - loss: 0.0982 - val_accuracy: 0.8825 - val_loss: 0.3187


Epoch 9/20
30/30  0s 11ms/step - accuracy: 0.9740 - loss: 0.0876 - val_accuracy: 0.8820 - val_loss: 0.3343


Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9791 - loss: 0.0752 - val_accuracy: 0.8747 - val_loss: 0.3821


Epoch 11/20
30/30  0s 12ms/step - accuracy: 0.9842 - loss: 0.0657 - val_accuracy: 0.8817 - val_loss: 0.3622


Epoch 12/20
30/30  0s 11ms/step - accuracy: 0.9872 - loss: 0.0546 - val_accuracy: 0.8787 - val_loss: 0.3831


Epoch 13/20
30/30  0s 12ms/step - accuracy: 0.9922 - loss: 0.0419 - val_accuracy: 0.8767 - val_loss: 0.4007


Epoch 14/20
30/30  0s 11ms/step - accuracy: 0.9940 - loss: 0.0363 - val_accuracy: 0.8721 - val_loss: 0.4465


Epoch 15/20
30/30  0s 11ms/step - accuracy: 0.9948 - loss: 0.0299 - val_accuracy: 0.8733 - val_loss: 0.4469


Epoch 16/20
30/30  0s 11ms/step - accuracy: 0.9967 - loss: 0.0259 - val_accuracy: 0.8611 - val_loss: 0.5083


Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9980 - loss: 0.0214 - val_accuracy: 0.8701 - val_loss: 0.4964


Epoch 18/20
30/30  0s 11ms/step - accuracy: 0.9981 - loss: 0.0179 - val_accuracy: 0.8739 - val_loss: 0.5164


Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9981 - loss: 0.0162 - val_accuracy: 0.8719 - val_loss: 0.5349


Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9996 - loss: 0.0103 - val_accuracy: 0.8715 - val_loss: 0.5610


Epoch 1/20
30/30  1s 22ms/step - accuracy: 0.6900 - loss: 0.2225 - val_accuracy: 0.8460 - val_loss: 0.1460


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8730 - loss: 0.1249 - val_accuracy: 0.8688 - val_loss: 0.1087


Epoch 3/20
30/30  0s 12ms/step - accuracy: 0.9041 - loss: 0.0889 - val_accuracy: 0.8721 - val_loss: 0.0981


Epoch 4/20
30/30  0s 11ms/step - accuracy: 0.9187 - loss: 0.0719 - val_accuracy: 0.8880 - val_loss: 0.0870


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.9369 - loss: 0.0593 - val_accuracy: 0.8865 - val_loss: 0.0849


Epoch 6/20
30/30  0s 11ms/step - accuracy: 0.9481 - loss: 0.0493 - val_accuracy: 0.8874 - val_loss: 0.0841


Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9582 - loss: 0.0431 - val_accuracy: 0.8741 - val_loss: 0.0926


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9609 - loss: 0.0385 - val_accuracy: 0.8855 - val_loss: 0.0832


Epoch 9/20
30/30  0s 11ms/step - accuracy: 0.9665 - loss: 0.0350 - val_accuracy: 0.8831 - val_loss: 0.0844


Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9733 - loss: 0.0297 - val_accuracy: 0.8718 - val_loss: 0.0965


Epoch 11/20
30/30  0s 11ms/step - accuracy: 0.9755 - loss: 0.0279 - val_accuracy: 0.8797 - val_loss: 0.0882


Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9810 - loss: 0.0235 - val_accuracy: 0.8792 - val_loss: 0.0875


Epoch 13/20
30/30  0s 11ms/step - accuracy: 0.9816 - loss: 0.0219 - val_accuracy: 0.8716 - val_loss: 0.0944


Epoch 14/20
30/30  0s 11ms/step - accuracy: 0.9846 - loss: 0.0202 - val_accuracy: 0.8784 - val_loss: 0.0903


Epoch 15/20
30/30  0s 11ms/step - accuracy: 0.9854 - loss: 0.0184 - val_accuracy: 0.8756 - val_loss: 0.0918

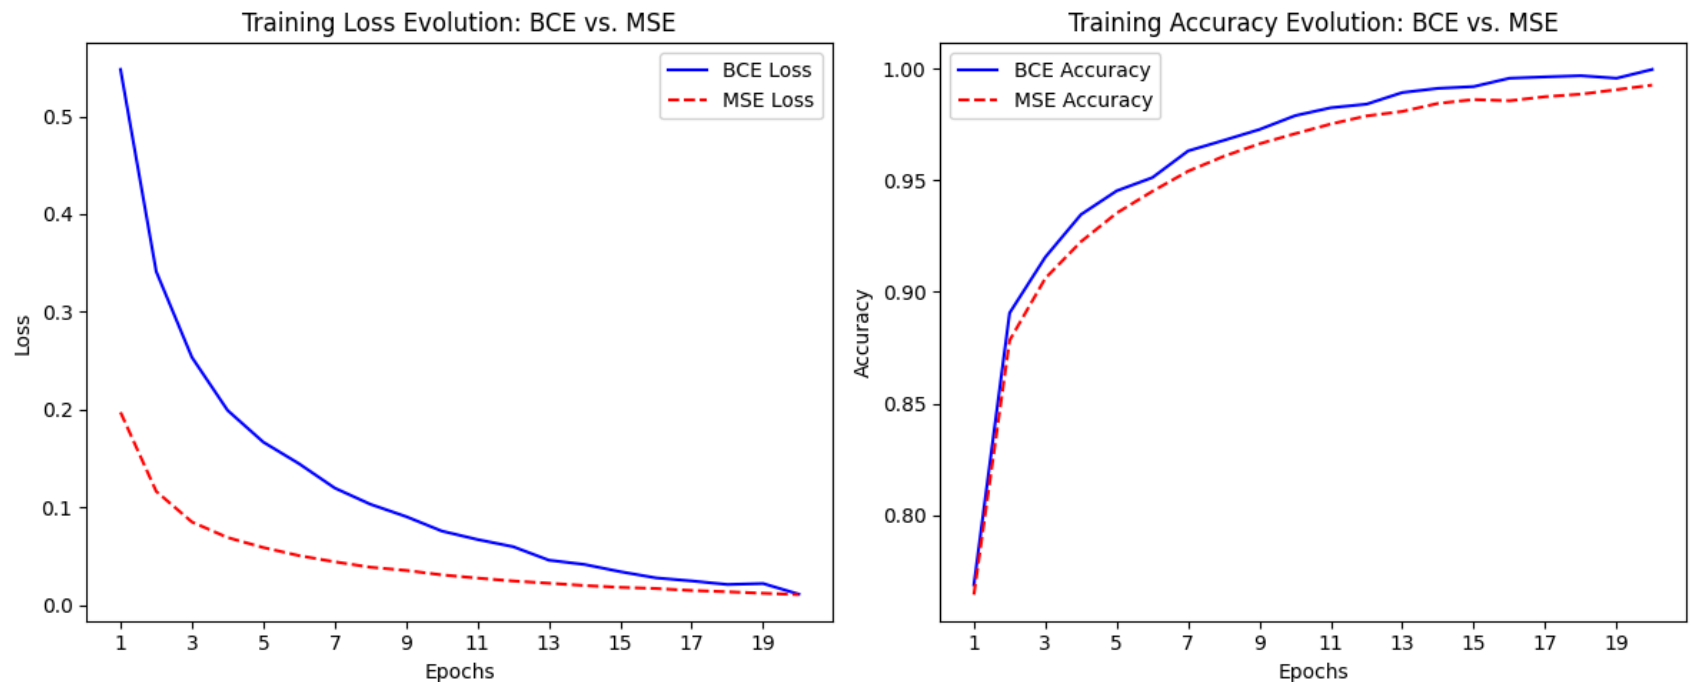
Epoch 16/20
30/30  0s 11ms/step - accuracy: 0.9870 - loss: 0.0161 - val_accuracy: 0.8745 - val_loss: 0.0931

Epoch 17/20
30/30  0s 11ms/step - accuracy: 0.9895 - loss: 0.0134 - val_accuracy: 0.8760 - val_loss: 0.0937

Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9917 - loss: 0.0114 - val_accuracy: 0.8739 - val_loss: 0.0954

Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9921 - loss: 0.0104 - val_accuracy: 0.8749 - val_loss: 0.0963

Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9929 - loss: 0.0097 - val_accuracy: 0.8754 - val_loss: 0.0972



Step 6: Experimenting with Activation Functions (ReLU vs. Tanh)

```
In [25]: # Train a model with ReLU activation
relu_model = build_model(activation="relu")
relu_history = relu_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                             validation_data=(x_val, y_val), verbose=1)

# Train a model with Tanh activation
tanh_model = build_model(activation="tanh")
tanh_history = tanh_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                              validation_data=(x_val, y_val), verbose=1)

# Plot Training Evolution for Loss and Accuracy over 20 Epochs
plt.figure(figsize=(12, 5))

# Plot Loss Evolution
plt.subplot(1, 2, 1)
plt.plot(range(1, 21), relu_history.history["loss"], label="ReLU Loss", color="blue")
plt.plot(range(1, 21), tanh_history.history["loss"], label="Tanh Loss", color="red", linestyle="--")
```

```

plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss Evolution: ReLU vs. Tanh")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


# Plot Accuracy Evolution
plt.subplot(1, 2, 2)
plt.plot(range(1, 21), relu_history.history["accuracy"], label="ReLU Accuracy", color="blue")
plt.plot(range(1, 21), tanh_history.history["accuracy"], label="Tanh Accuracy", color="red", linestyle="--")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Training Accuracy Evolution: ReLU vs. Tanh")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


plt.tight_layout()
plt.show()


```


C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/20
30/30  2s 27ms/step - accuracy: 0.6894 - loss: 0.6287 - val_accuracy: 0.8039 - val_loss: 0.4755


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8701 - loss: 0.3985 - val_accuracy: 0.8772 - val_loss: 0.3408


Epoch 3/20
30/30  0s 11ms/step - accuracy: 0.9080 - loss: 0.2854 - val_accuracy: 0.8717 - val_loss: 0.3165


Epoch 4/20
30/30  0s 12ms/step - accuracy: 0.9222 - loss: 0.2302 - val_accuracy: 0.8855 - val_loss: 0.2863


Epoch 5/20
30/30  0s 11ms/step - accuracy: 0.9391 - loss: 0.1871 - val_accuracy: 0.8876 - val_loss: 0.2786


Epoch 6/20
30/30  0s 11ms/step - accuracy: 0.9429 - loss: 0.1645 - val_accuracy: 0.8862 - val_loss: 0.2787


Epoch 7/20
30/30  0s 11ms/step - accuracy: 0.9548 - loss: 0.1423 - val_accuracy: 0.8831 - val_loss: 0.2910


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9615 - loss: 0.1237 - val_accuracy: 0.8831 - val_loss: 0.3089


Epoch 9/20
30/30  0s 11ms/step - accuracy: 0.9685 - loss: 0.1108 - val_accuracy: 0.8792 - val_loss: 0.3147


Epoch 10/20
30/30  0s 11ms/step - accuracy: 0.9738 - loss: 0.0932 - val_accuracy: 0.8749 - val_loss: 0.3396


Epoch 11/20
30/30  0s 12ms/step - accuracy: 0.9778 - loss: 0.0800 - val_accuracy: 0.8800 - val_loss: 0.3352


Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9809 - loss: 0.0707 - val_accuracy: 0.8744 - val_loss: 0.3673


Epoch 13/20
30/30  0s 11ms/step - accuracy: 0.9849 - loss: 0.0618 - val_accuracy: 0.8704 - val_loss: 0.3893


Epoch 14/20
30/30  0s 11ms/step - accuracy: 0.9882 - loss: 0.0544 - val_accuracy: 0.8719 - val_loss: 0.4018


Epoch 15/20
30/30  0s 12ms/step - accuracy: 0.9912 - loss: 0.0450 - val_accuracy: 0.8737 - val_loss: 0.4167


Epoch 16/20
30/30  0s 12ms/step - accuracy: 0.9939 - loss: 0.0381 - val_accuracy: 0.8711 - val_loss: 0.4426


Epoch 17/20
30/30  0s 11ms/step - accuracy: 0.9938 - loss: 0.0348 - val_accuracy: 0.8725 - val_loss: 0.4705


Epoch 18/20
30/30  0s 11ms/step - accuracy: 0.9961 - loss: 0.0288 - val_accuracy: 0.8711 - val_loss: 0.4929


Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9960 - loss: 0.0260 - val_accuracy: 0.8693 - val_loss: 0.5047


Epoch 20/20
30/30  0s 11ms/step - accuracy: 0.9976 - loss: 0.0218 - val_accuracy: 0.8676 - val_loss: 0.5350


Epoch 1/20
30/30  1s 23ms/step - accuracy: 0.7076 - loss: 0.5911 - val_accuracy: 0.8407 - val_loss: 0.4154


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8932 - loss: 0.3358 - val_accuracy: 0.8717 - val_loss: 0.3223


Epoch 3/20
30/30  0s 12ms/step - accuracy: 0.9240 - loss: 0.2352 - val_accuracy: 0.8807 - val_loss: 0.2877


Epoch 4/20
30/30  0s 11ms/step - accuracy: 0.9433 - loss: 0.1772 - val_accuracy: 0.8860 - val_loss: 0.2726


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.9528 - loss: 0.1462 - val_accuracy: 0.8845 - val_loss: 0.2845


Epoch 6/20
30/30  0s 12ms/step - accuracy: 0.9689 - loss: 0.1064 - val_accuracy: 0.8832 - val_loss: 0.3049


Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9729 - loss: 0.0926 - val_accuracy: 0.8822 - val_loss: 0.3292


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9801 - loss: 0.0721 - val_accuracy: 0.8664 - val_loss: 0.3985


Epoch 9/20
30/30  0s 12ms/step - accuracy: 0.9845 - loss: 0.0595 - val_accuracy: 0.8554 - val_loss: 0.4753


Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9823 - loss: 0.0617 - val_accuracy: 0.8688 - val_loss: 0.4753


Epoch 11/20
30/30  0s 12ms/step - accuracy: 0.9911 - loss: 0.0382 - val_accuracy: 0.8710 - val_loss: 0.4579


Epoch 12/20
30/30  0s 11ms/step - accuracy: 0.9946 - loss: 0.0291 - val_accuracy: 0.8666 - val_loss: 0.4963


Epoch 13/20
30/30  0s 11ms/step - accuracy: 0.9961 - loss: 0.0233 - val_accuracy: 0.8654 - val_loss: 0.5313


Epoch 14/20
30/30  0s 12ms/step - accuracy: 0.9981 - loss: 0.0155 - val_accuracy: 0.8660 - val_loss: 0.5560


Epoch 15/20
30/30  0s 12ms/step - accuracy: 0.9933 - loss: 0.0278 - val_accuracy: 0.8674 - val_loss: 0.5796

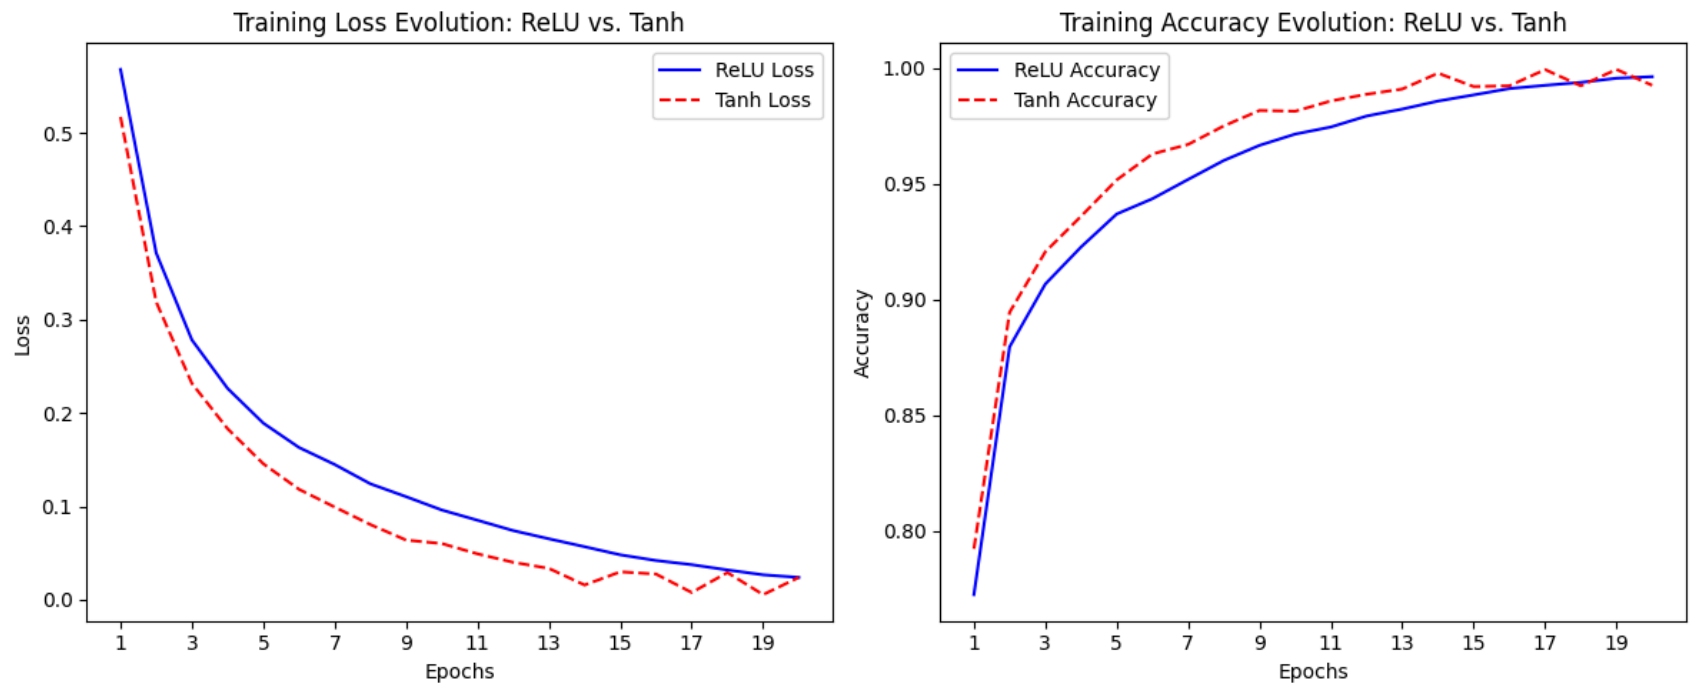
Epoch 16/20
30/30  0s 13ms/step - accuracy: 0.9973 - loss: 0.0137 - val_accuracy: 0.8648 - val_loss: 0.6032

Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9995 - loss: 0.0087 - val_accuracy: 0.8634 - val_loss: 0.6358

Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9932 - loss: 0.0279 - val_accuracy: 0.8660 - val_loss: 0.6406

Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9995 - loss: 0.0055 - val_accuracy: 0.8647 - val_loss: 0.6682

Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9965 - loss: 0.0131 - val_accuracy: 0.8650 - val_loss: 0.6811



Step 7: Applying Regularization Techniques

```
In [26]: # Train a model WITHOUT Dropout (Standard Model)
no_dropout_model = build_model()
no_dropout_history = no_dropout_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                                         validation_data=(x_val, y_val), verbose=1)

# Train a model WITH Dropout Regularization
dropout_model = build_model_with_dropout()
dropout_history = dropout_model.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512,
                                   validation_data=(x_val, y_val), verbose=1)

# Plot Training Evolution for Loss and Accuracy over 20 Epochs
plt.figure(figsize=(12, 5))

# Plot Loss Evolution
plt.subplot(1, 2, 1)
plt.plot(range(1, 21), no_dropout_history.history["loss"], label="No Dropout Loss", color="blue")
plt.plot(range(1, 21), dropout_history.history["loss"], label="Dropout Loss", color="red", linestyle="--")
```

```

plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss Evolution: No Dropout vs. Dropout")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


# Plot Accuracy Evolution
plt.subplot(1, 2, 2)
plt.plot(range(1, 21), no_dropout_history.history["accuracy"], label="No Dropout Accuracy", color="blue")
plt.plot(range(1, 21), dropout_history.history["accuracy"], label="Dropout Accuracy", color="red", linestyle="--")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Training Accuracy Evolution: No Dropout vs. Dropout")
plt.xticks(range(1, 21, 2)) # Show all 20 epochs clearly
plt.legend()


plt.tight_layout()
plt.show()


```


C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/20
30/30  1s 26ms/step - accuracy: 0.6978 - loss: 0.6022 - val_accuracy: 0.8643 - val_loss: 0.3925


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8908 - loss: 0.3397 - val_accuracy: 0.8823 - val_loss: 0.3155


Epoch 3/20
30/30  0s 11ms/step - accuracy: 0.9250 - loss: 0.2436 - val_accuracy: 0.8402 - val_loss: 0.3687


Epoch 4/20
30/30  0s 12ms/step - accuracy: 0.9297 - loss: 0.2079 - val_accuracy: 0.8816 - val_loss: 0.2921


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.9495 - loss: 0.1604 - val_accuracy: 0.8885 - val_loss: 0.2762


Epoch 6/20
30/30  0s 12ms/step - accuracy: 0.9626 - loss: 0.1306 - val_accuracy: 0.8838 - val_loss: 0.2897


Epoch 7/20
30/30  0s 11ms/step - accuracy: 0.9664 - loss: 0.1146 - val_accuracy: 0.8735 - val_loss: 0.3235


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9730 - loss: 0.0971 - val_accuracy: 0.8832 - val_loss: 0.3063


Epoch 9/20
30/30  0s 12ms/step - accuracy: 0.9789 - loss: 0.0793 - val_accuracy: 0.8809 - val_loss: 0.3400


Epoch 10/20
30/30  0s 11ms/step - accuracy: 0.9801 - loss: 0.0728 - val_accuracy: 0.8808 - val_loss: 0.3563


Epoch 11/20
30/30  0s 11ms/step - accuracy: 0.9853 - loss: 0.0616 - val_accuracy: 0.8816 - val_loss: 0.3560


Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9920 - loss: 0.0459 - val_accuracy: 0.8774 - val_loss: 0.3770


Epoch 13/20
30/30  0s 11ms/step - accuracy: 0.9929 - loss: 0.0412 - val_accuracy: 0.8785 - val_loss: 0.4008


Epoch 14/20
30/30  0s 12ms/step - accuracy: 0.9947 - loss: 0.0338 - val_accuracy: 0.8659 - val_loss: 0.4552


Epoch 15/20
30/30  0s 11ms/step - accuracy: 0.9925 - loss: 0.0361 - val_accuracy: 0.8763 - val_loss: 0.4420


Epoch 16/20
30/30  0s 12ms/step - accuracy: 0.9965 - loss: 0.0225 - val_accuracy: 0.8747 - val_loss: 0.4636




















Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9980 - loss: 0.0183 - val_accuracy: 0.8752 - val_loss: 0.4866

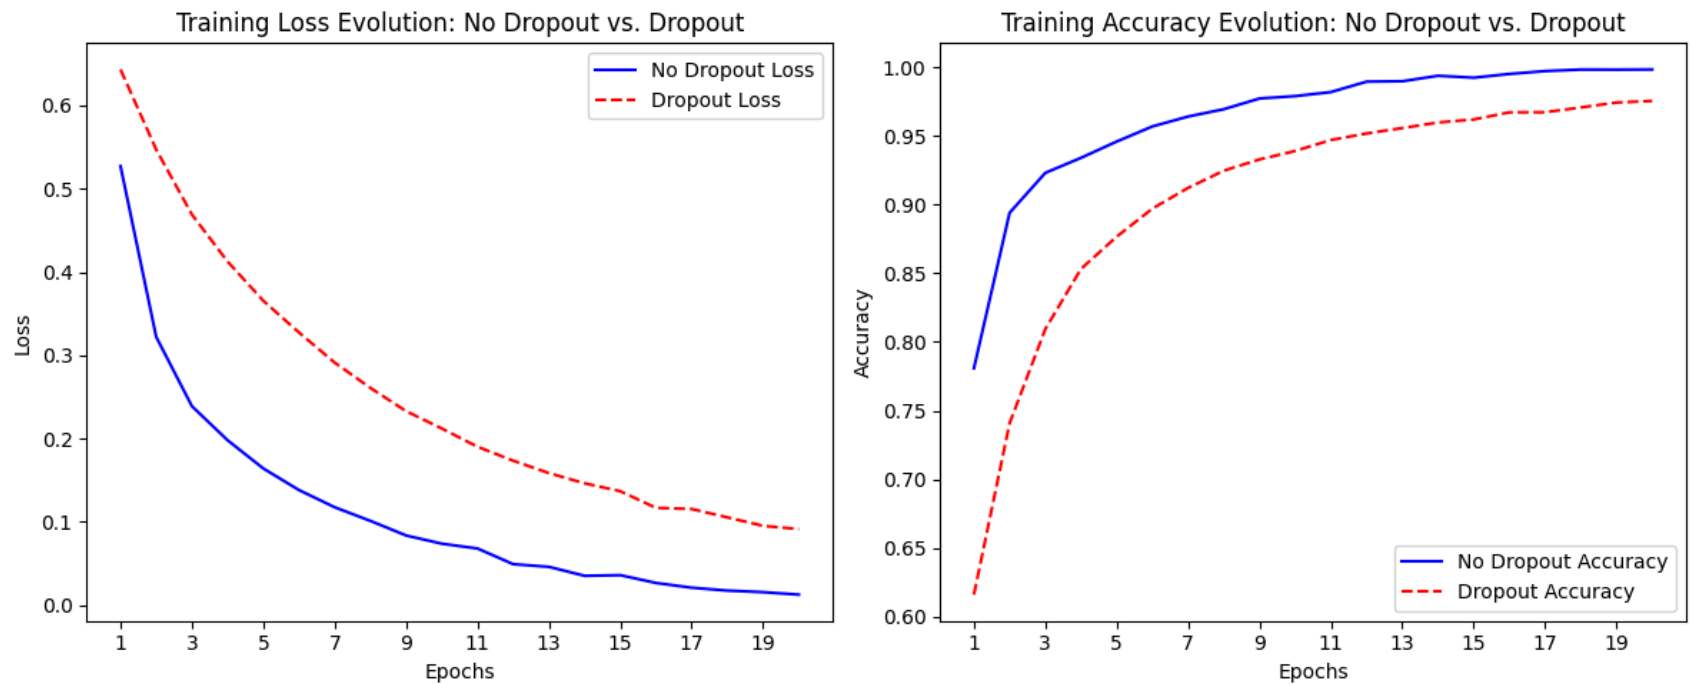
Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9991 - loss: 0.0158 - val_accuracy: 0.8733 - val_loss: 0.5126

Epoch 19/20
30/30  0s 11ms/step - accuracy: 0.9992 - loss: 0.0132 - val_accuracy: 0.8723 - val_loss: 0.5402

Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9987 - loss: 0.0116 - val_accuracy: 0.8696 - val_loss: 0.5681

Epoch 1/20
30/30  1s 23ms/step - accuracy: 0.5702 - loss: 0.6688 - val_accuracy: 0.8274 - val_loss: 0.5467

Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.7239 - loss: 0.5633 - val_accuracy: 0.8641 - val_loss: 0.4496
Epoch 3/20
30/30  0s 12ms/step - accuracy: 0.8023 - loss: 0.4779 - val_accuracy: 0.8772 - val_loss: 0.3844
Epoch 4/20
30/30  0s 12ms/step - accuracy: 0.8464 - loss: 0.4237 - val_accuracy: 0.8822 - val_loss: 0.3407
Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.8760 - loss: 0.3694 - val_accuracy: 0.8830 - val_loss: 0.3142
Epoch 6/20
30/30  0s 12ms/step - accuracy: 0.8953 - loss: 0.3305 - val_accuracy: 0.8899 - val_loss: 0.2872
Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9132 - loss: 0.2916 - val_accuracy: 0.8819 - val_loss: 0.3085
Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9253 - loss: 0.2621 - val_accuracy: 0.8849 - val_loss: 0.3013
Epoch 9/20
30/30  0s 12ms/step - accuracy: 0.9348 - loss: 0.2313 - val_accuracy: 0.8870 - val_loss: 0.2892
Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9400 - loss: 0.2109 - val_accuracy: 0.8893 - val_loss: 0.2943
Epoch 11/20
30/30  0s 12ms/step - accuracy: 0.9486 - loss: 0.1885 - val_accuracy: 0.8874 - val_loss: 0.2964
Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9502 - loss: 0.1778 - val_accuracy: 0.8878 - val_loss: 0.3162
Epoch 13/20
30/30  0s 12ms/step - accuracy: 0.9566 - loss: 0.1588 - val_accuracy: 0.8870 - val_loss: 0.3287
Epoch 14/20
30/30  0s 12ms/step - accuracy: 0.9600 - loss: 0.1490 - val_accuracy: 0.8855 - val_loss: 0.3464
Epoch 15/20
30/30  0s 12ms/step - accuracy: 0.9605 - loss: 0.1416 - val_accuracy: 0.8852 - val_loss: 0.3785
Epoch 16/20
30/30  0s 12ms/step - accuracy: 0.9673 - loss: 0.1150 - val_accuracy: 0.8837 - val_loss: 0.4131
Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9687 - loss: 0.1100 - val_accuracy: 0.8824 - val_loss: 0.4121
Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9691 - loss: 0.1045 - val_accuracy: 0.8826 - val_loss: 0.4372
Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9759 - loss: 0.0939 - val_accuracy: 0.8824 - val_loss: 0.4618
Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9755 - loss: 0.0927 - val_accuracy: 0.8789 - val_loss: 0.5138



Step 8: Dropout and L2 Regularization

```
In [27]: # Train a model with Dropout and L2 Regularization
def build_model_with_dropout_and_l2():
    model = keras.Sequential([
        layers.Dense(16, activation="relu", kernel_regularizer=keras.regularizers.l2(0.001), input_shape=(10000,)),
        layers.Dropout(0.3),
        layers.Dense(16, activation="relu", kernel_regularizer=keras.regularizers.l2(0.001)),
        layers.Dropout(0.3),
        layers.Dense(1, activation="sigmoid")
    ])
    model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
    return model

dropout_l2_model = build_model_with_dropout_and_l2()
dropout_l2_history = dropout_l2_model.fit(
    partial_x_train, partial_y_train, epochs=20, batch_size=512,
    validation_data=(x_val, y_val), verbose=1
)
```

```

# Plot training and validation loss
import matplotlib.pyplot as plt

history = dropout_l2_history.history

epochs = range(1, 21)
plt.figure(figsize=(12, 5))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(epochs, history["loss"], "bo", label="Training Loss")
plt.plot(epochs, history["val_loss"], "b", label="Validation Loss")
plt.title("Training and Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()


# Plot accuracy
plt.subplot(1, 2, 2)
plt.plot(epochs, history["accuracy"], "bo", label="Training Accuracy")
plt.plot(epochs, history["val_accuracy"], "b", label="Validation Accuracy")
plt.title("Training and Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()


plt.show()


```


C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/20
30/30  1s 26ms/step - accuracy: 0.5880 - loss: 0.7011 - val_accuracy: 0.8533 - val_loss: 0.4986


Epoch 2/20
30/30  0s 12ms/step - accuracy: 0.8141 - loss: 0.5039 - val_accuracy: 0.8526 - val_loss: 0.4188


Epoch 3/20
30/30  0s 12ms/step - accuracy: 0.8729 - loss: 0.3959 - val_accuracy: 0.8748 - val_loss: 0.3620


Epoch 4/20
30/30  0s 12ms/step - accuracy: 0.8969 - loss: 0.3411 - val_accuracy: 0.8877 - val_loss: 0.3348


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.9157 - loss: 0.2953 - val_accuracy: 0.8875 - val_loss: 0.3299


Epoch 6/20
30/30  0s 13ms/step - accuracy: 0.9256 - loss: 0.2660 - val_accuracy: 0.8780 - val_loss: 0.3555


Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9439 - loss: 0.2326 - val_accuracy: 0.8834 - val_loss: 0.3454


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9497 - loss: 0.2152 - val_accuracy: 0.8798 - val_loss: 0.3467


Epoch 9/20
30/30  0s 12ms/step - accuracy: 0.9580 - loss: 0.1998 - val_accuracy: 0.8851 - val_loss: 0.3503


Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9612 - loss: 0.1877 - val_accuracy: 0.8750 - val_loss: 0.4256


Epoch 11/20
30/30  0s 12ms/step - accuracy: 0.9642 - loss: 0.1789 - val_accuracy: 0.8797 - val_loss: 0.3686


Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9678 - loss: 0.1679 - val_accuracy: 0.8805 - val_loss: 0.3839


Epoch 13/20
30/30  0s 12ms/step - accuracy: 0.9735 - loss: 0.1552 - val_accuracy: 0.8805 - val_loss: 0.4061


Epoch 14/20
30/30  0s 12ms/step - accuracy: 0.9731 - loss: 0.1512 - val_accuracy: 0.8786 - val_loss: 0.4107


Epoch 15/20
30/30  0s 12ms/step - accuracy: 0.9740 - loss: 0.1488 - val_accuracy: 0.8780 - val_loss: 0.4125

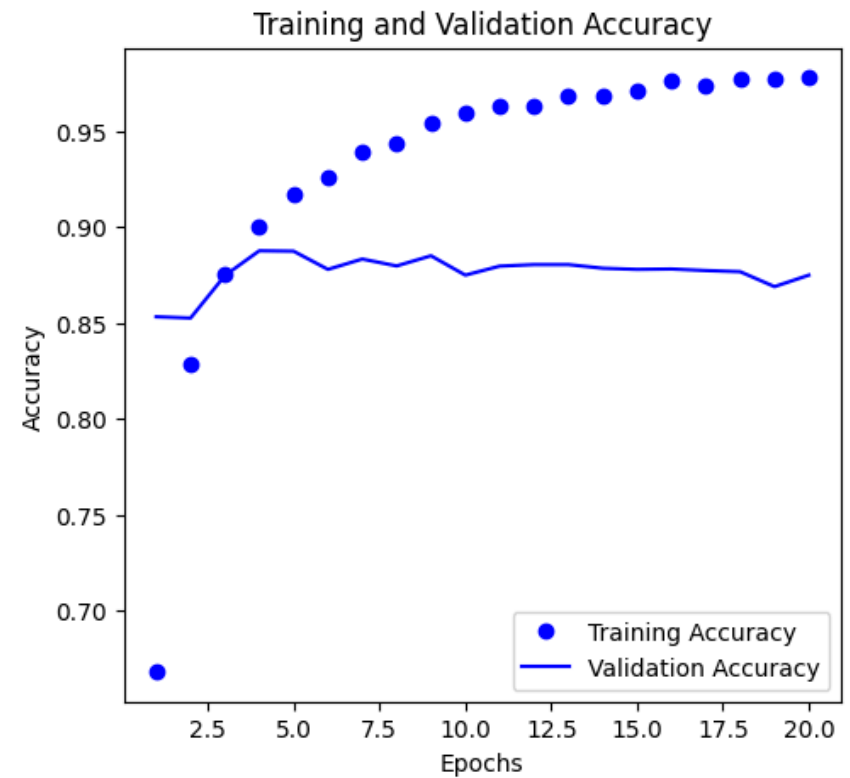
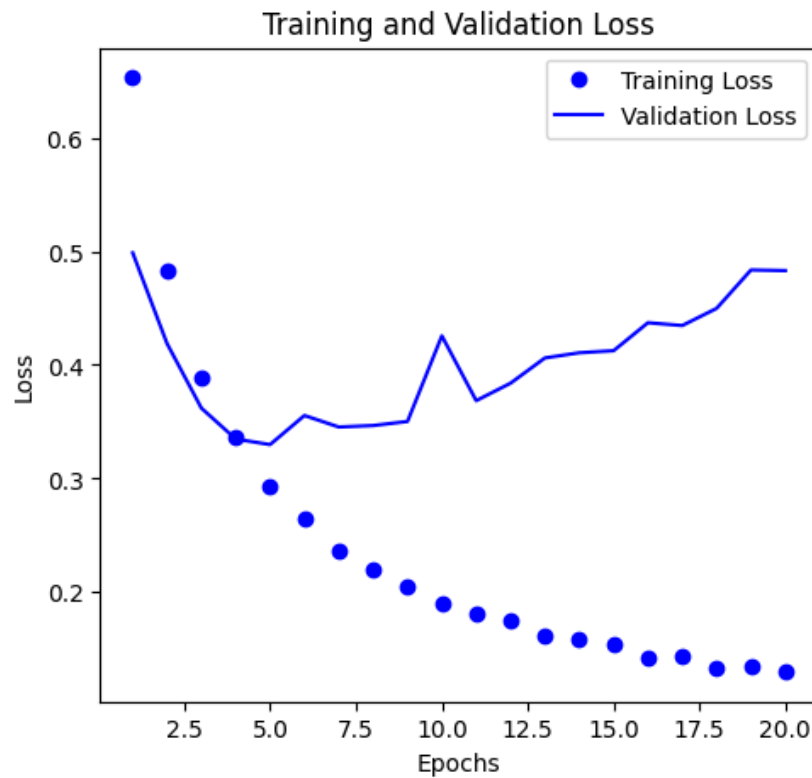
Epoch 16/20
30/30  0s 12ms/step - accuracy: 0.9788 - loss: 0.1396 - val_accuracy: 0.8782 - val_loss: 0.4371

Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9782 - loss: 0.1338 - val_accuracy: 0.8774 - val_loss: 0.4347

Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9818 - loss: 0.1280 - val_accuracy: 0.8768 - val_loss: 0.4498

Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9814 - loss: 0.1250 - val_accuracy: 0.8690 - val_loss: 0.4836

Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9795 - loss: 0.1263 - val_accuracy: 0.8750 - val_loss: 0.4830



```
In [28]: # Train a model with L2 regularization and adjusted dropout
def build_model_with_l2_and_dropout():
    model = keras.Sequential([
        layers.Dense(16, activation="relu", input_shape=(10000,)), kernel_regularizer=keras.regularizers.l2(0.002)),
        layers.Dropout(0.4),
        layers.Dense(16, activation="relu", kernel_regularizer=keras.regularizers.l2(0.002)),
        layers.Dropout(0.4),
        layers.Dense(1, activation="sigmoid")
    ])
    model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
    return model

# Train the model with L2 regularization and adjusted dropout
model_with_l2_and_dropout = build_model_with_l2_and_dropout()
history_with_l2_and_dropout = model_with_l2_and_dropout.fit(partial_x_train, partial_y_train, epochs=20, batch_size=100,
                                                            validation_data=(x_val, y_val), verbose=1)
```



```

# Plot the results
import matplotlib.pyplot as plt

# Plot training & validation accuracy values
plt.figure(figsize=(12, 6))
plt.plot(history_with_l2_and_dropout.history['accuracy'], label='Training Accuracy')
plt.plot(history_with_l2_and_dropout.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy with L2 Regularization and Dropout (0.4)')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

# Plot training & validation loss values
plt.figure(figsize=(12, 6))
plt.plot(history_with_l2_and_dropout.history['loss'], label='Training Loss')
plt.plot(history_with_l2_and_dropout.history['val_loss'], label='Validation Loss')
plt.title('Model Loss with L2 Regularization and Dropout (0.4)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')

plt.show()


```


C:\Users\zhanguu\AppData\Local\miniconda3\envs\tf_env2\lib\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```


super().__init__(activity_regularizer=activity_regularizer, **kwargs)


```


Epoch 1/20
30/30  1s 25ms/step - accuracy: 0.5770 - loss: 0.7590 - val_accuracy: 0.8447 - val_loss: 0.6106


Epoch 2/20
30/30  0s 11ms/step - accuracy: 0.7599 - loss: 0.6051 - val_accuracy: 0.8706 - val_loss: 0.4889


Epoch 3/20
30/30  0s 12ms/step - accuracy: 0.8247 - loss: 0.5051 - val_accuracy: 0.8700 - val_loss: 0.4412


Epoch 4/20
30/30  0s 12ms/step - accuracy: 0.8653 - loss: 0.4446 - val_accuracy: 0.8846 - val_loss: 0.3889


Epoch 5/20
30/30  0s 12ms/step - accuracy: 0.8885 - loss: 0.3945 - val_accuracy: 0.8826 - val_loss: 0.3786


Epoch 6/20
30/30  0s 12ms/step - accuracy: 0.9019 - loss: 0.3584 - val_accuracy: 0.8840 - val_loss: 0.3571


Epoch 7/20
30/30  0s 12ms/step - accuracy: 0.9201 - loss: 0.3298 - val_accuracy: 0.8848 - val_loss: 0.3553


Epoch 8/20
30/30  0s 12ms/step - accuracy: 0.9274 - loss: 0.3123 - val_accuracy: 0.8850 - val_loss: 0.3624


Epoch 9/20
30/30  0s 12ms/step - accuracy: 0.9317 - loss: 0.2969 - val_accuracy: 0.8862 - val_loss: 0.3663


Epoch 10/20
30/30  0s 12ms/step - accuracy: 0.9339 - loss: 0.2799 - val_accuracy: 0.8825 - val_loss: 0.3635


Epoch 11/20
30/30  0s 13ms/step - accuracy: 0.9432 - loss: 0.2648 - val_accuracy: 0.8816 - val_loss: 0.3740


Epoch 12/20
30/30  0s 12ms/step - accuracy: 0.9441 - loss: 0.2585 - val_accuracy: 0.8803 - val_loss: 0.3752


Epoch 13/20
30/30  0s 12ms/step - accuracy: 0.9516 - loss: 0.2449 - val_accuracy: 0.8805 - val_loss: 0.3809


Epoch 14/20
30/30  0s 12ms/step - accuracy: 0.9522 - loss: 0.2322 - val_accuracy: 0.8833 - val_loss: 0.3959


Epoch 15/20
30/30  0s 12ms/step - accuracy: 0.9571 - loss: 0.2239 - val_accuracy: 0.8740 - val_loss: 0.4689

Epoch 16/20
30/30  0s 12ms/step - accuracy: 0.9514 - loss: 0.2275 - val_accuracy: 0.8782 - val_loss: 0.4412

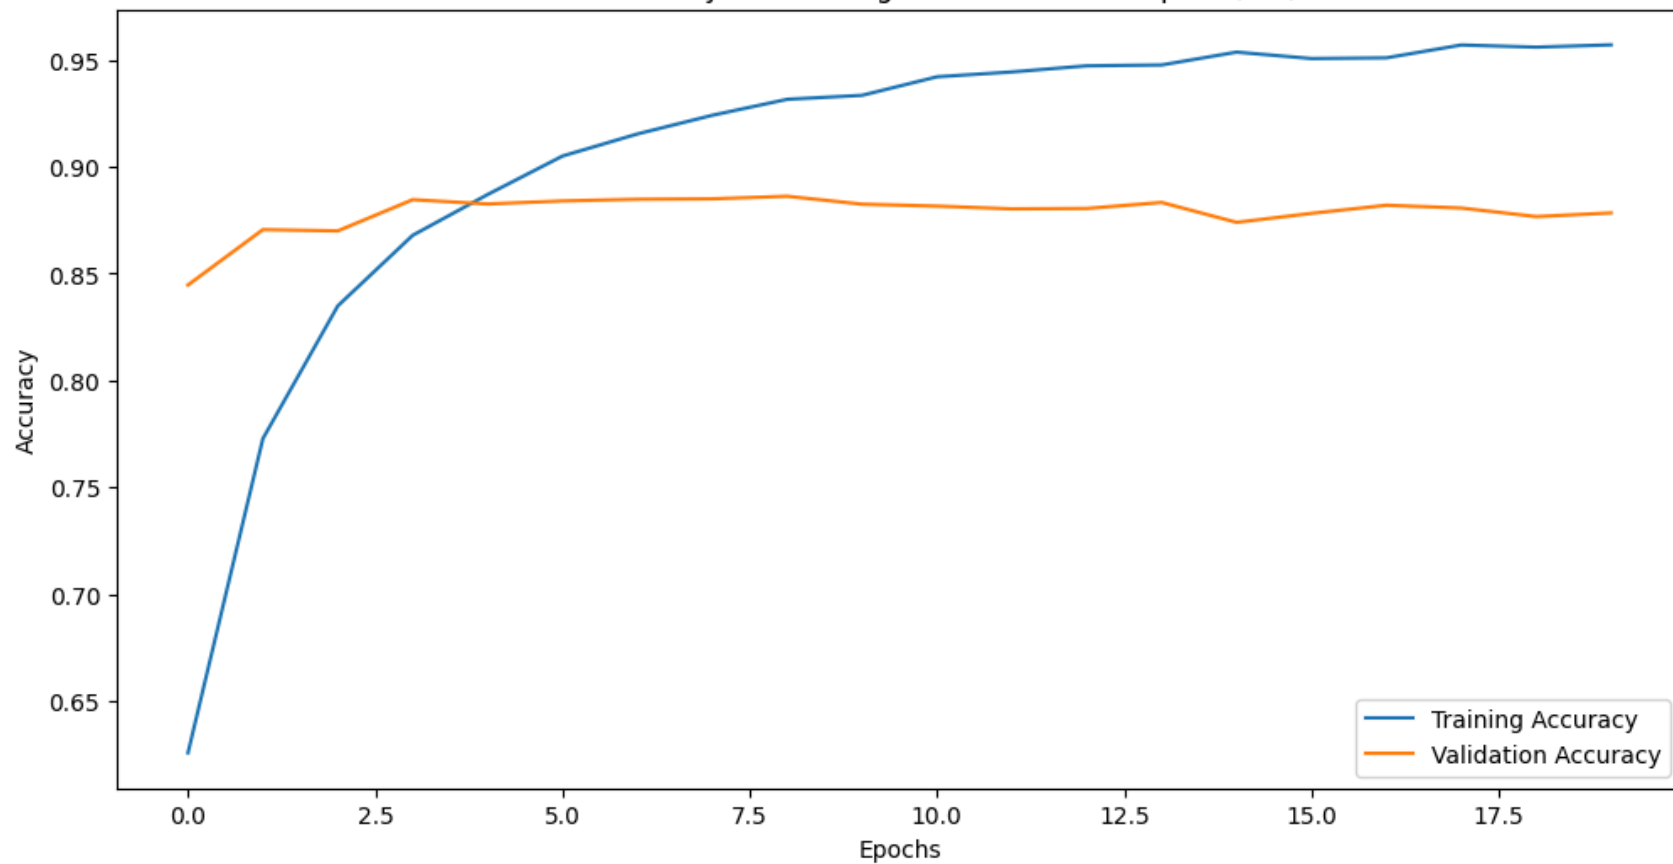
Epoch 17/20
30/30  0s 12ms/step - accuracy: 0.9554 - loss: 0.2233 - val_accuracy: 0.8820 - val_loss: 0.4189

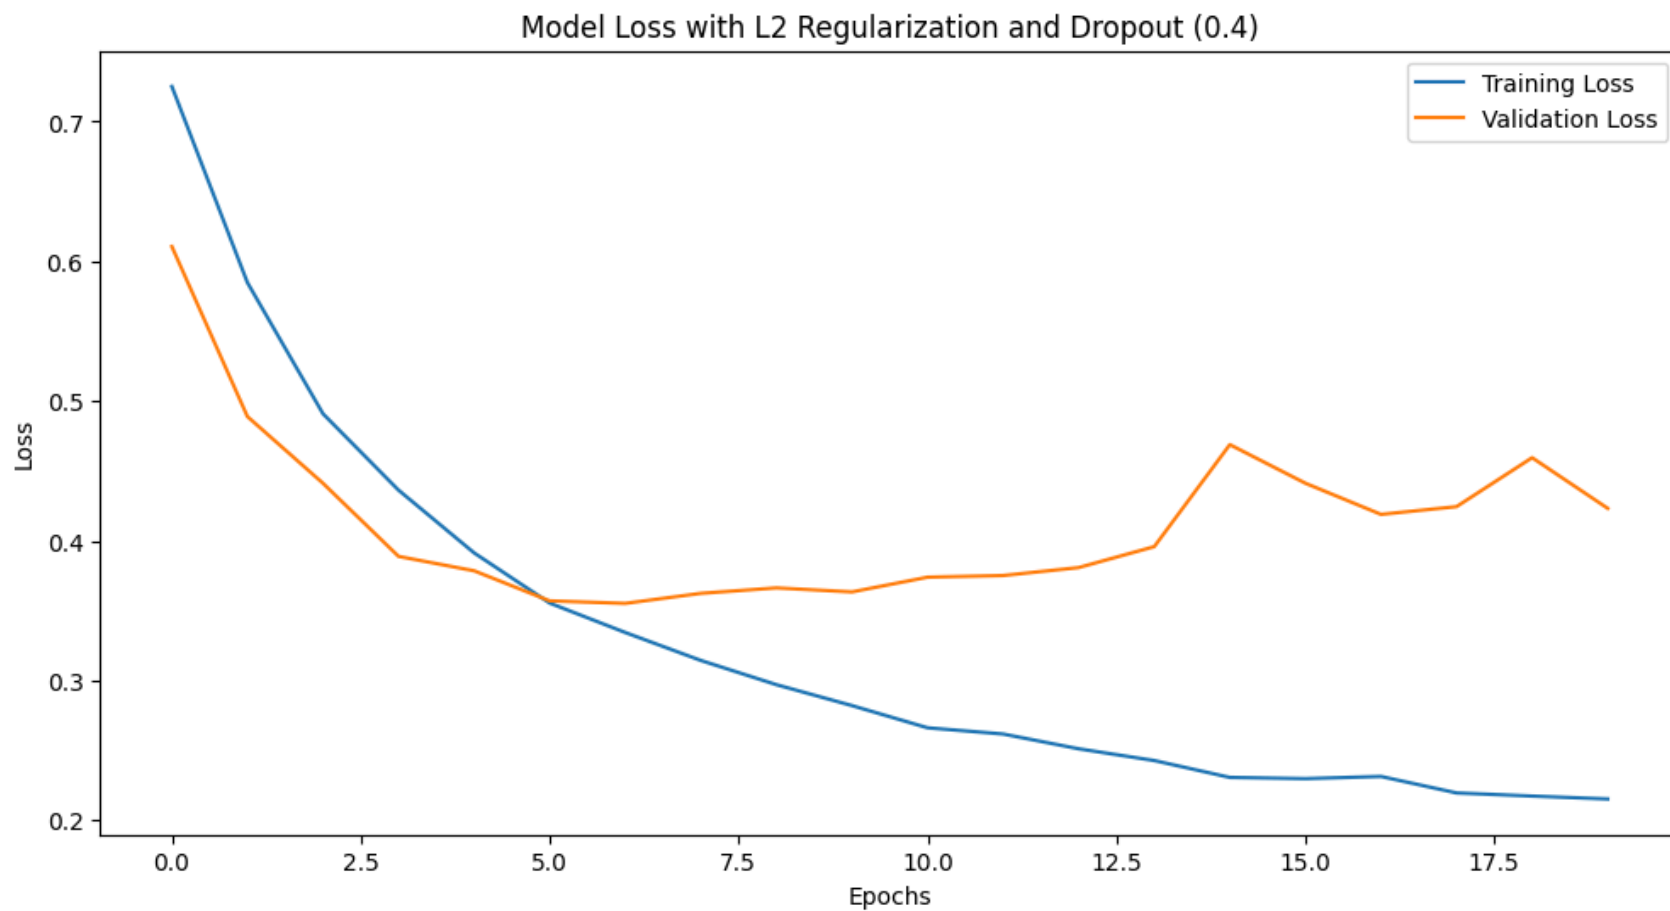
Epoch 18/20
30/30  0s 12ms/step - accuracy: 0.9612 - loss: 0.2091 - val_accuracy: 0.8807 - val_loss: 0.4245

Epoch 19/20
30/30  0s 12ms/step - accuracy: 0.9605 - loss: 0.2096 - val_accuracy: 0.8767 - val_loss: 0.4595

Epoch 20/20
30/30  0s 12ms/step - accuracy: 0.9612 - loss: 0.2081 - val_accuracy: 0.8784 - val_loss: 0.4233

Model Accuracy with L2 Regularization and Dropout (0.4)





高速下载